

# CS 380: Artificial Intelligence

---

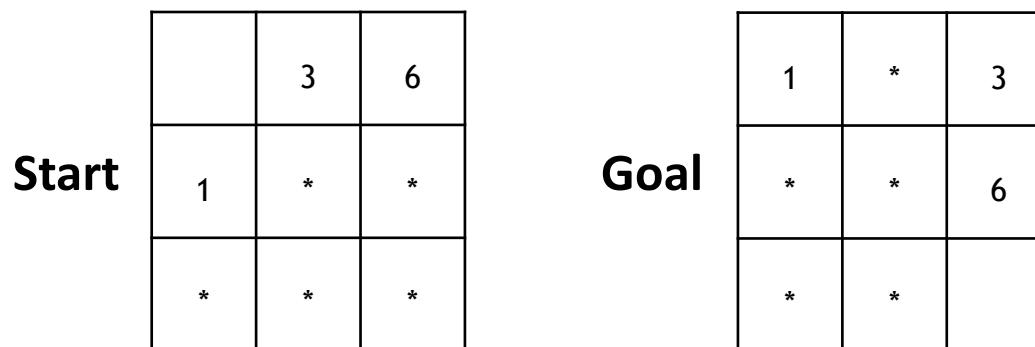
## Lecture 5: Local Search & CSPs

---

# A General Method for Building Heuristics

---

- **Pattern Databases** – Store the exact optimal solution cost for every possible subproblem instance.
  - You care about some features of the problem, but not all
  - In this example, all the possible locations of  $[1,3,6,blank]$
  - Identify the instance in the DB to get the  $h$  value



# Recall: Problem Solving

---

- Idea: represent the problem we want to solve as:
  - State space
  - Actions
  - Goal check
  - Heuristic / evaluation function
- If this is possible, we can use search methods to solve any problem:
  - BFS, DFS, ID, Greedy Search, A\*

# Systematic vs. Local Search

---

- All the algorithms we have seen so far are “systematic”:
  - They will systematically explore ALL the states in the search space if needed
    - each of them just explores it in a different order
  - If the search space is finite, they are all complete
- Local Search:
  - Instead of systematically exploring the state space...
  - Start with one state (even a bad one), and iteratively improve it by exploring neighbor states (“local”)
  - Search space is the set of all states visited
  - *Iterative improvement*

## Iterative improvement algorithms

In many optimization problems, **path** is irrelevant;  
the goal state itself is the solution

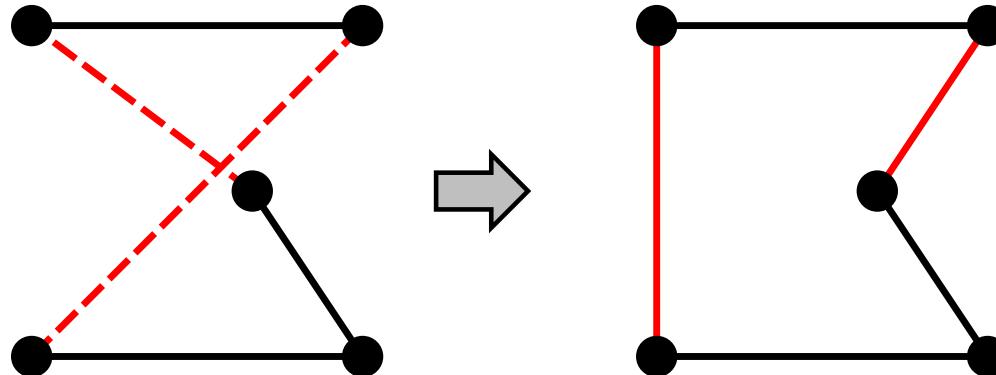
Then state space = set of “complete” configurations;  
find **optimal** configuration, e.g., TSP  
or, find configuration satisfying constraints, e.g., timetable

In such cases, can use **iterative improvement** algorithms;  
keep a single “current” state, try to improve it

Constant space, suitable for online as well as offline search

## Example: Travelling Salesperson Problem

Start with any complete tour, perform pairwise exchanges

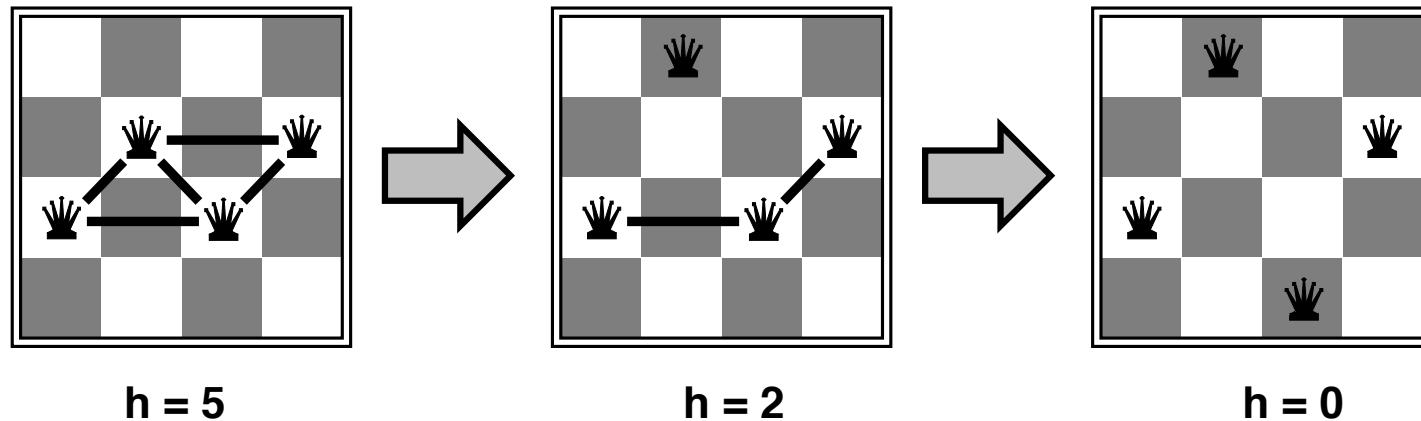


Variants of this approach get within 1% of optimal very quickly with thousands of cities

## Example: $n$ -queens

Put  $n$  queens on an  $n \times n$  board with no two queens on the same row, column, or diagonal

Move a queen to reduce number of conflicts



Almost always solves  $n$ -queens problems almost instantaneously for very large  $n$ , e.g.,  $n=1\text{million}$

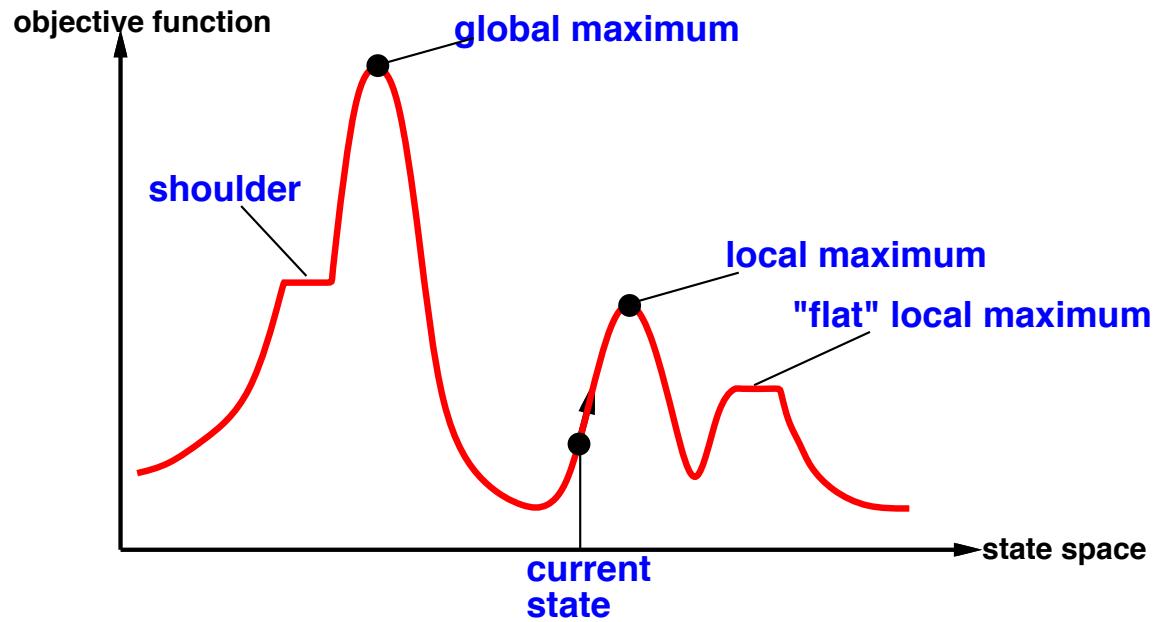
## Hill-climbing (or gradient ascent/descent)

“Like climbing Everest in thick fog with amnesia”

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
    inputs: problem, a problem
    local variables: current, a node
                    neighbor, a node
    current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])
    loop do
        neighbor  $\leftarrow$  a highest-valued successor of current
        if VALUE[neighbor]  $\leq$  VALUE[current] then return STATE[current]
        current  $\leftarrow$  neighbor
    end
```

## Hill-climbing contd.

Useful to consider state space landscape



Random-restart hill climbing overcomes local maxima—trivially complete

Random sideways moves 😊 escape from shoulders 😓 loop on flat maxima

# Simulated Annealing

---

- “Annealing” = heating and then controlled cooling (of metals to strengthen the material)
- At a fixed temperature  $T$ , the state occupation probability reaches the Boltzmann distribution.

$$p(x) = \alpha e^{\frac{E(x)}{kT}}$$

- If  $T$  is decreased slowly enough you are guaranteed to reach the optimal state (if we need optimal!?)
- Visualization:
  - [https://www.youtube.com/watch?v=iaq\\_Fpr4KZc](https://www.youtube.com/watch?v=iaq_Fpr4KZc)

# Simulated annealing

Idea: escape local maxima by allowing some “bad” moves  
**but gradually decrease their size and frequency**

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
    inputs: problem, a problem
            schedule, a mapping from time to “temperature”
    local variables: current, a node
                      next, a node
                      T, a “temperature” controlling prob. of downward steps

    current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])
    for t  $\leftarrow$  1 to  $\infty$  do
        T  $\leftarrow$  schedule[t]
        if T = 0 then return current
        next  $\leftarrow$  a randomly selected successor of current
         $\Delta E \leftarrow$  VALUE[next] – VALUE[current]
        if  $\Delta E > 0$  then current  $\leftarrow$  next
        else current  $\leftarrow$  next only with probability  $e^{\Delta E/T}$ 
```

## Properties of simulated annealing

At fixed “temperature”  $T$ , state occupation probability reaches Boltzman distribution

$$p(x) = \alpha e^{\frac{E(x)}{kT}}$$

$T$  decreased slowly enough  $\implies$  always reach best state  $x^*$   
because  $e^{\frac{E(x^*)}{kT}} / e^{\frac{E(x)}{kT}} = e^{\frac{E(x^*) - E(x)}{kT}} \gg 1$  for small  $T$

Is this necessarily an interesting guarantee??

Devised by Metropolis et al., 1953, for physical process modelling

Widely used in VLSI layout, airline scheduling, etc.

## Local beam search

Idea: keep  $k$  states instead of 1; choose top  $k$  of all their successors

Not the same as  $k$  searches run in parallel!

Searches that find good states recruit other searches to join them

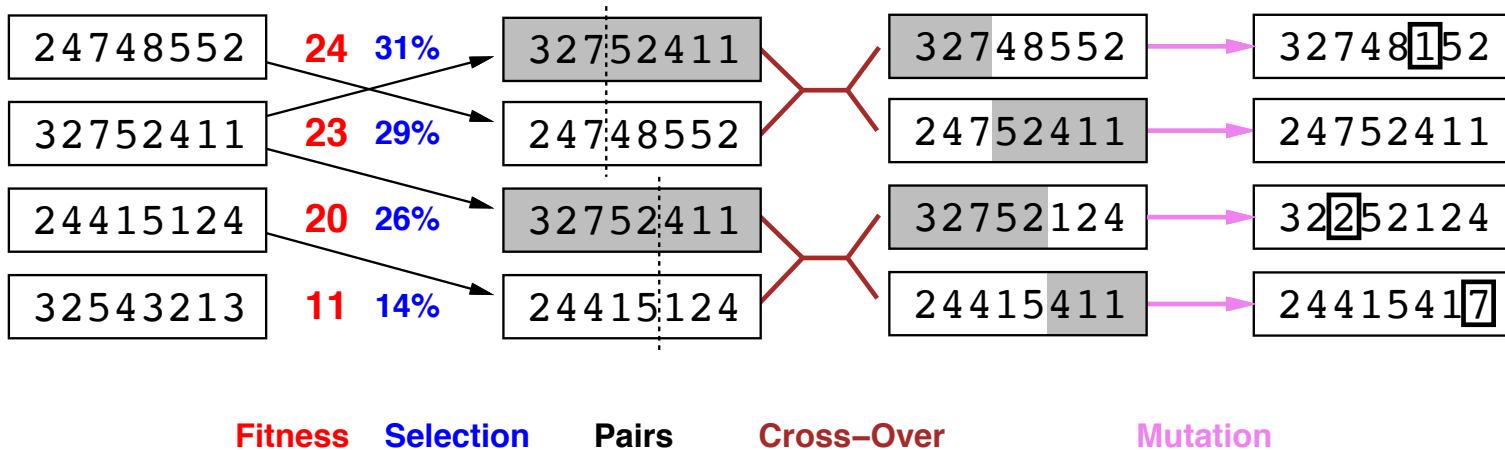
Problem: quite often, all  $k$  states end up on same local hill

Idea: choose  $k$  successors randomly, biased towards good ones

Observe the close analogy to natural selection!

# Genetic algorithms

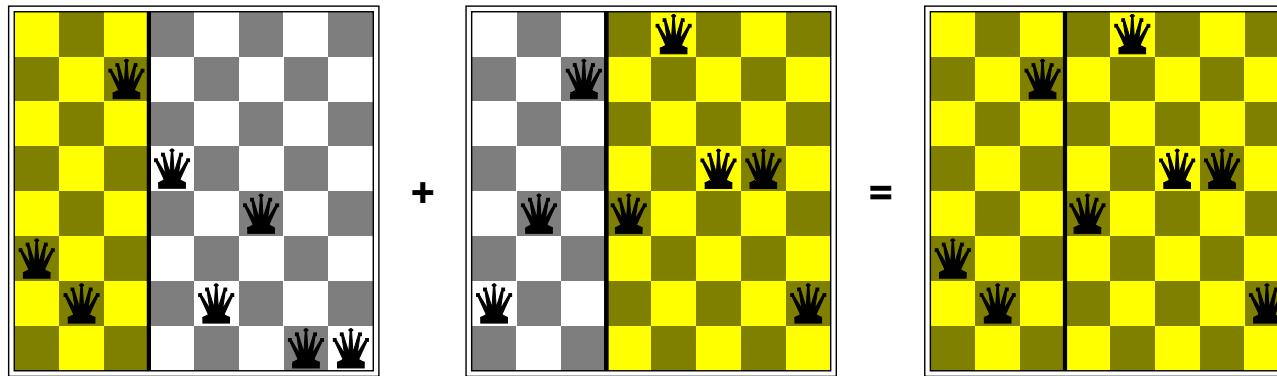
= stochastic local beam search + generate successors from **pairs** of states



## Genetic algorithms contd.

GAs require states encoded as strings ([GPs](#) use programs)

Crossover helps **iff substrings are meaningful components**



GAs  $\neq$  evolution: e.g., real genes encode replication machinery!

# Sample Visualization

---

- Greedy search, local search, simulated annealing...

<https://www.youtube.com/watch?v=SC5CX8drAtU>

## Continuous state spaces

Suppose we want to site three airports in Romania:

- 6-D state space defined by  $(x_1, y_2)$ ,  $(x_2, y_2)$ ,  $(x_3, y_3)$
- objective function  $f(x_1, y_2, x_2, y_2, x_3, y_3) =$   
sum of squared distances from each city to nearest airport

Discretization methods turn continuous space into discrete space,  
e.g., empirical gradient considers  $\pm\delta$  change in each coordinate

Gradient methods compute

$$\nabla f = \left( \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial y_1}, \frac{\partial f}{\partial x_2}, \frac{\partial f}{\partial y_2}, \frac{\partial f}{\partial x_3}, \frac{\partial f}{\partial y_3} \right)$$

to increase/reduce  $f$ , e.g., by  $\mathbf{x} \leftarrow \mathbf{x} + \alpha \nabla f(\mathbf{x})$

Sometimes can solve for  $\nabla f(\mathbf{x}) = 0$  exactly (e.g., with one city).

Newton–Raphson (1664, 1690) iterates  $\mathbf{x} \leftarrow \mathbf{x} - \mathbf{H}_f^{-1}(\mathbf{x}) \nabla f(\mathbf{x})$   
to solve  $\nabla f(\mathbf{x}) = 0$ , where  $\mathbf{H}_{ij} = \partial^2 f / \partial x_i \partial x_j$

## Constraint satisfaction problems (CSPs)

Standard search problem:

state is a “black box”—any old data structure  
that supports goal test, eval, successor

CSP:

state is defined by *variables*  $V_i$  with *values* from *domain*  $D_i$

goal test is a set of *constraints* specifying  
allowable combinations of values for subsets of variables

Simple example of a *formal representation language*

Allows useful *general-purpose* algorithms with more power  
than standard search algorithms

## Example: 4-Queens as a CSP

Assume one queen in each column. Which row does each one go in?

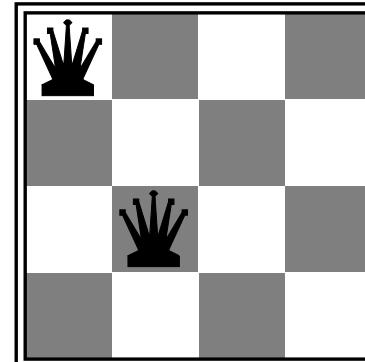
Variables  $Q_1, Q_2, Q_3, Q_4$

Domains  $D_i = \{1, 2, 3, 4\}$

Constraints

$Q_i \neq Q_j$  (cannot be in same row)

$|Q_i - Q_j| \neq |i - j|$  (or same diagonal)



$$Q_1 = 1 \quad Q_2 = 3$$

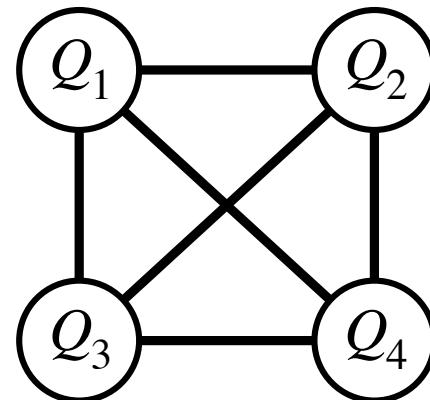
Translate each constraint into set of allowable values for its variables

E.g., values for  $(Q_1, Q_2)$  are  $(1, 3)$   $(1, 4)$   $(2, 4)$   $(3, 1)$   $(4, 1)$   $(4, 2)$

## Constraint graph

*Binary CSP*: each constraint relates at most two variables

*Constraint graph*: nodes are variables, arcs show constraints



## Example: Cryptarithmetic

### Variables

$D \ E \ M \ N \ O \ R \ S \ Y$

### Domains

$\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$

$$\begin{array}{r} \text{S} \ \text{E} \ \text{N} \ \text{D} \\ + \ \text{M} \ \text{O} \ \text{R} \ \text{E} \\ \hline \text{M} \ \text{O} \ \text{N} \ \text{E} \ \text{Y} \end{array}$$

### Constraints

$M \neq 0, S \neq 0$  (*unary constraints*)

$Y = D + E$  or  $Y = D + E - 10$ , etc.

$D \neq E, D \neq M, D \neq N$ , etc.

## Example: Map coloring

Color a map so that no adjacent countries have the same color

### Variables

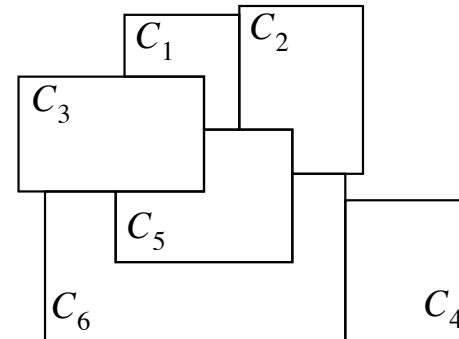
Countries  $C_i$

### Domains

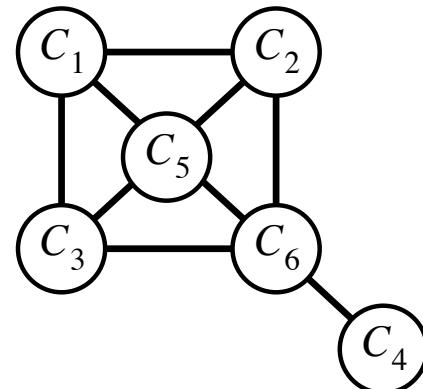
$\{Red, Blue, Green\}$

### Constraints

$C_1 \neq C_2, C_1 \neq C_5$ , etc.



Constraint graph:



## Real-world CSPs

Assignment problems

e.g., who teaches what class

Timetabling problems

e.g., which class is offered when and where?

Hardware configuration

Spreadsheets

Transportation scheduling

Factory scheduling

Floorplanning

Notice that many real-world problems involve real-valued variables

## Applying standard search

Let's start with the straightforward, dumb approach, then fix it

States are defined by the values assigned so far

Initial state: all variables unassigned

Operators: assign a value to an unassigned variable

Goal test: all variables assigned, no constraints violated

Notice that this is the same for all CSPs!

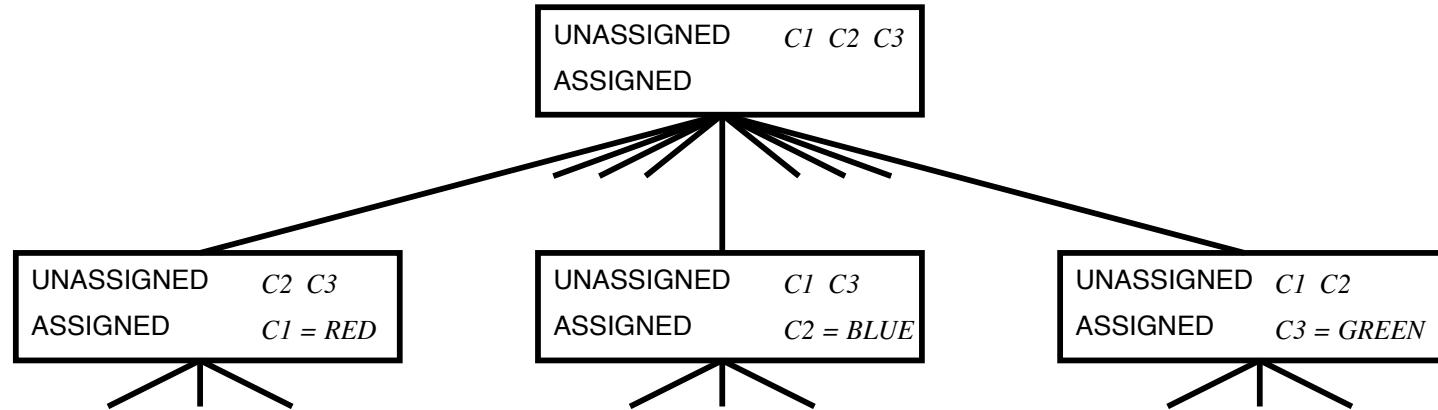
# Implementation

CSP state keeps track of which variables have values so far  
Each variable has a domain and a current value

```
datatype CSP-STATE
  components: UNASSIGNED, a list of variables not yet assigned
              ASSIGNED, a list of variables that have values
datatype CSP-VAR
  components: NAME, for i/o purposes
              DOMAIN, a list of possible values
              VALUE, current value (if any)
```

Constraints can be represented  
explicitly as sets of allowable values, or  
implicitly by a function that tests for satisfaction of the constraint

## Standard search applied to map-coloring



## Complexity of the dumb approach

Max. depth of space  $m = ??$

Depth of solution state  $d = ??$

Search algorithm to use??

Branching factor  $b = ??$

This can be improved dramatically by noting the following:

- 1) Order of assignment is irrelevant, hence many paths are equivalent
- 2) Adding assignments cannot correct a violated constraint

## Complexity of the dumb approach

Max. depth of space  $m = ?? n$  (number of variables)

Depth of solution state  $d = ?? n$  (all vars assigned)

Search algorithm to use?? depth-first

Branching factor  $b = ?? \sum_i |D_i|$  (at top of tree)       $|D_i|$  = size of Domain  $i$

This can be improved dramatically by noting the following:

- 1) Order of assignment is irrelevant so many paths are equivalent
- 2) Adding assignments cannot correct a violated constraint

## Backtracking search

Use depth-first search, but

- 1) fix the order of assignment,  $\Rightarrow b = |D_i|$   
(can be done in the SUCCESSORS function)
- 2) check for constraint violations

The constraint violation check can be implemented in two ways:

- 1) modify SUCCESSORS to assign only values that  
are allowed, given the values already assigned
- or 2) check constraints are satisfied before expanding a state

Backtracking search is the basic uninformed algorithm for CSPs

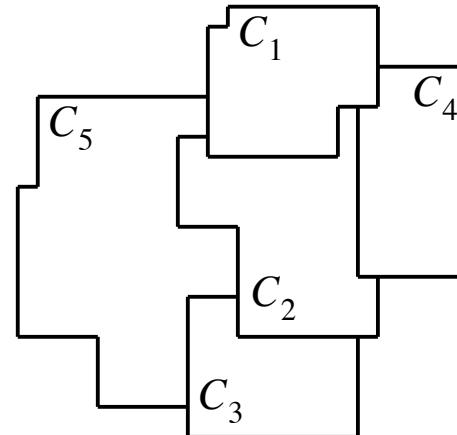
Can solve  $n$ -queens for  $n \approx 15$

## Forward checking

Idea: Keep track of remaining legal values for unassigned variables  
Terminate search when any variable has no legal values

Simplified map-coloring example:

	RED	BLUE	GREEN
$C_1$			
$C_2$			
$C_3$			
$C_4$			
$C_5$			



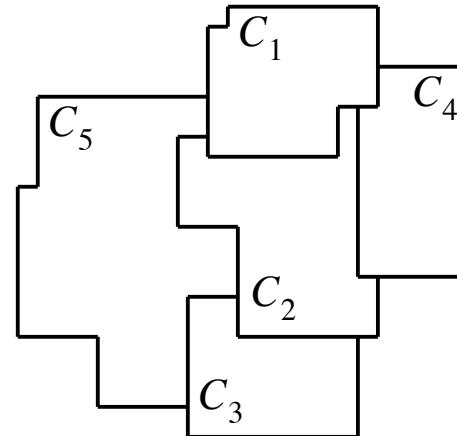
Can solve  $n$ -queens up to  $n \approx 30$

## Forward checking

Idea: Keep track of remaining legal values for unassigned variables  
Terminate search when any variable has no legal values

Simplified map-coloring example:

	RED	BLUE	GREEN
$C_1$	✓		
$C_2$	✗		
$C_3$			
$C_4$	✗		
$C_5$	✗		



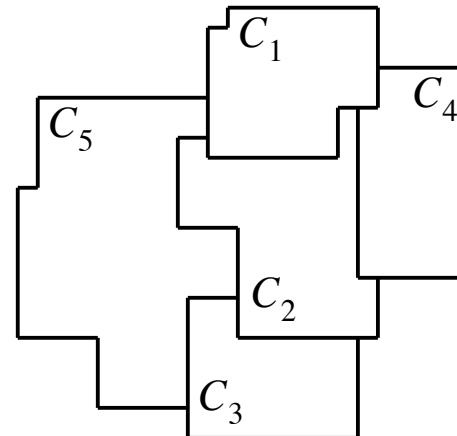
Can solve  $n$ -queens up to  $n \approx 30$

## Forward checking

Idea: Keep track of remaining legal values for unassigned variables  
Terminate search when any variable has no legal values

Simplified map-coloring example:

	RED	BLUE	GREEN
$C_1$	✓		
$C_2$	✗	✓	
$C_3$		✗	
$C_4$	✗	✗	
$C_5$	✗	✗	



Can solve  $n$ -queens up to  $n \approx 30$

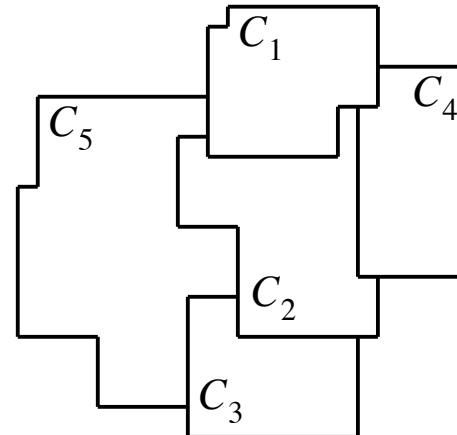
## Forward checking

Idea: Keep track of remaining legal values for unassigned variables  
Terminate search when any variable has no legal values

Simplified map-coloring example:

	RED	BLUE	GREEN
$C_1$	✓		
$C_2$	✗	✓	
$C_3$		✗	✓
$C_4$	✗	✗	
$C_5$	✗	✗	✗

← problem!



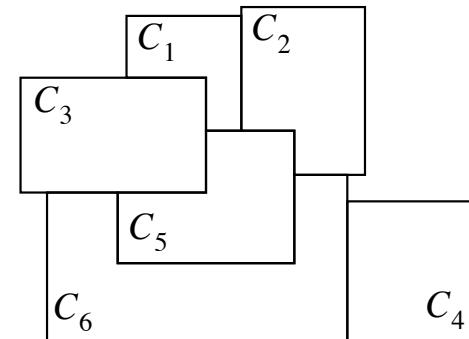
Can solve  $n$ -queens up to  $n \approx 30$

# Heuristics for CSPs

More intelligent decisions on  
which value to choose for each variable  
which variable to assign next

Given  $C_1 = Red, C_2 = Green$ , choose  $C_3 = ??$

Given  $C_1 = Red, C_2 = Green$ , what next??



Can solve  $n$ -queens for  $n \approx 1000$

## Heuristics for CSPs

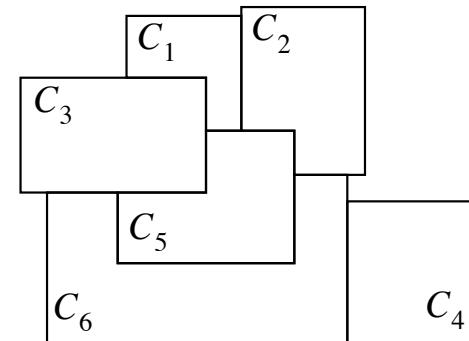
More intelligent decisions on  
which value to choose for each variable  
which variable to assign next

Given  $C_1 = \text{Red}$ ,  $C_2 = \text{Green}$ , choose  $C_3 = ??$

$C_3 = \text{Green}$ : least-constraining-value

Given  $C_1 = \text{Red}$ ,  $C_2 = \text{Green}$ , what next??

$C_5$ : most-constrained-variable



Can solve  $n$ -queens for  $n \approx 1000$

## Iterative algorithms for CSPs

Hill-climbing, simulated annealing typically work with “complete” states, i.e., all variables assigned

To apply to CSPs:

- allow states with unsatisfied constraints
- operators *reassign* variable values

Variable selection: randomly select any conflicted variable

*min-conflicts* heuristic:

- choose value that violates the fewest constraints
- i.e., hillclimb with  $h(n)$  = total number of violated constraints

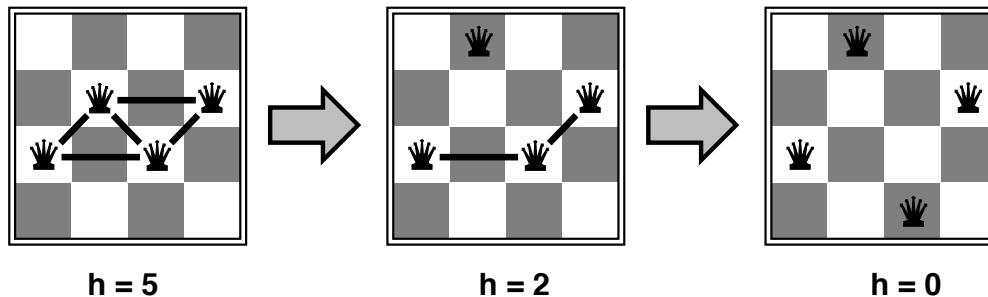
## Example: 4-Queens

States: 4 queens in 4 columns ( $4^4 = 256$  states)

Operators: move queen in column

Goal test: no attacks

Evaluation:  $h(n) = \text{number of attacks}$

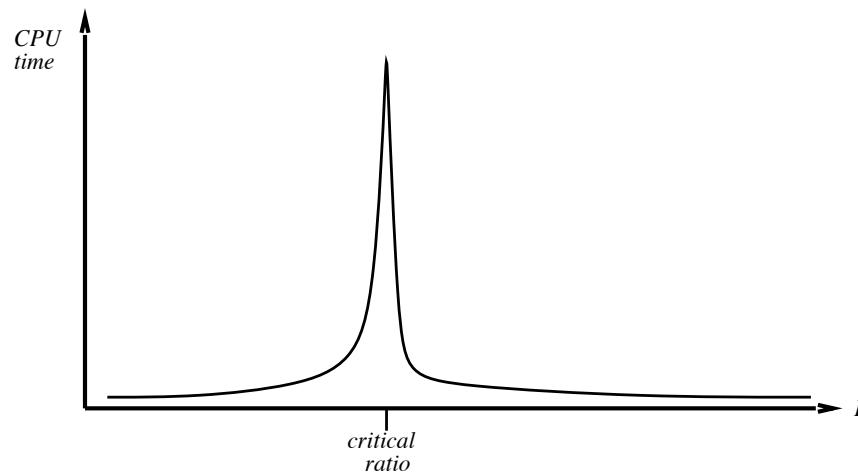


## Performance of min-conflicts

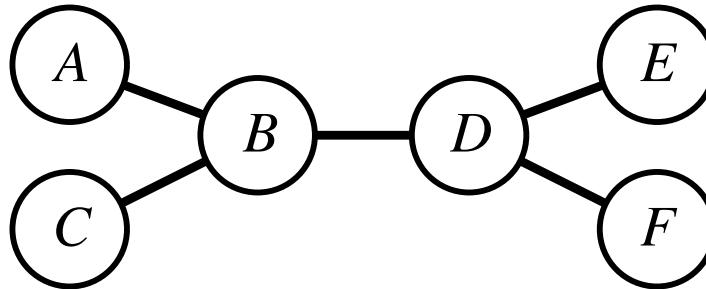
Given random initial state, can solve  $n$ -queens in almost constant time for arbitrary  $n$  with high probability (e.g.,  $n = 10,000,000$ )

The same appears to be true for any randomly-generated CSP  
except in a narrow range of the ratio

$$R = \frac{\text{number of constraints}}{\text{number of variables}}$$



## Tree-structured CSPs



Theorem: if the constraint graph has no loops, the CSP can be solved in  $O(n|D|^2)$  time

Compare to general CSPs, where worst-case time is  $O(|D|^n)$

This property also applies to logical and probabilistic reasoning:  
an important example of the relation between syntactic restrictions and complexity of reasoning.

## Algorithm for tree-structured CSPs

Basic step is called *filtering*:

$\text{FILTER}(V_i, V_j)$

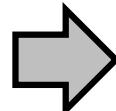
removes values of  $V_i$  that are inconsistent with ALL values of  $V_j$

Filtering example:



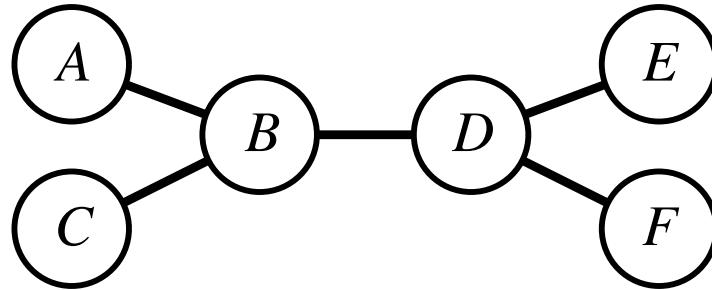
**allowed pairs:**

- < 1, 1 >
- < 3, 2 >
- < 3, 3 >

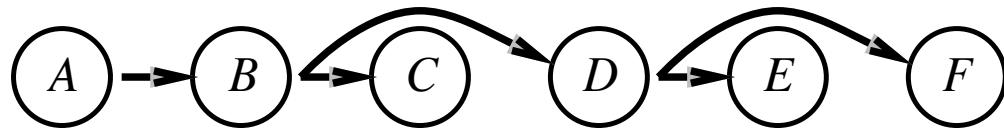


**remove 2 from  
domain of  $V_i$**

## Algorithm contd.



- 1) Order nodes breadth-first starting from any leaf:



- 2) For  $j = n$  to 1, apply FILTER( $V_i, V_j$ ) where  $V_i$  is a parent of  $V_j$
- 3) For  $j = 1$  to  $n$ , pick legal value for  $V_j$  given parent value

## Summary

CSPs are a special kind of problem:

states defined by values of a fixed set of variables

goal test defined by *constraints* on variable values

Backtracking = depth-first search with

- 1) fixed variable order
- 2) only legal successors

Forward checking prevents assignments that guarantee later failure

Variable ordering and value selection heuristics help significantly

Iterative min-conflicts is usually effective in practice

Tree-structured CSPs can always be solved very efficiently

# Coming up...

---

- We will start applying ideas of search to adversarial domains, where one player is playing another...