

CS 380: Artificial Intelligence

Lecture 12: Machine Learning

Some materials adapted from Russell & Norvig textbook slides: <http://aima.cs.berkeley.edu>

Summary so far:

- Rational Agents
- Problem Solving
 - Systematic Search
 - Uninformed
 - Informed
 - Local Search
 - Adversarial Search
- Logic and Knowledge Representation
 - Predicate Logic
 - First-order Logic
- Today: Machine Learning


What is Learning?

Machine Learning

- Computational methods for computers to exhibit specific forms of learning. For example:
 - Learning from Examples:
 - Supervised learning
 - Unsupervised learning
 - Reinforcement Learning
 - Learning from Observation (demonstration/imitation)

Examples

- **Supervised Learning:** learning to recognize writing




Examples

- **Supervised Learning:** image classification



Examples

- **Unsupervised Learning:** clustering observations into meaningful classes




Examples

- **Reinforcement Learning: learning to walk**



Examples

- **Learning from demonstration:** performing tasks that other agents (or humans) can do



Examples

- **Reinforcement Learning:**

- https://www.youtube.com/watch?v=hx_bgoTF7bs
- <https://www.youtube.com/watch?v=e27TUmMkOA0>

- **Learning from demonstration:**


- <https://www.youtube.com/watch?v=l9u3B6gv-RE>

Inductive Learning

- Inductive learning tries to infer general concepts from simple examples
- Let's try to learn something in its simplest form:
Learning a function from examples
 - **Desired:** f is a target function we want to learn
 - but we don't know what it is or how to compute it
 - **Given:** examples where an example = $\langle x, f(x) \rangle$
 - **Problem:** find a hypothesis h such that $h \approx f$


Inductive Learning

- Now, let's try to construct a hypothesis h such that $h \approx f$




Inductive Learning

- Now, let's try to construct a hypothesis h such that $h \approx f$




Inductive Learning

- Now, let's try to construct a hypothesis h such that $h \approx f$




Inductive Learning

- Now, let's try to construct a hypothesis h such that $h \approx f$




Inductive Learning

- Now, let's try to construct a hypothesis h such that $h \approx f$



Inductive Learning

- Now, let's try to construct a hypothesis h such that $h \approx f$



- Occam's razor:** The simplest explanation tends to be the correct/best(?) one.



William of Occam
(1287-1347)

Attribute-Based Representations

- Items (or examples) described by attribute values (Boolean, discrete, continuous, etc.)
- Example: Will I wait for a table at a restaurant?
 - What factors (attributes) would affect your decision?

Attribute-Based Representations

- Items (or examples) described by attribute values (Boolean, discrete, continuous, etc.)
- Example: Will I wait for a table at a restaurant?
 - Do I have an **Alternative**?
 - Is there a **Bar**?
 - Is it **Friday**?
 - Am I **Hungry**?
 - Are there **Patrons** in the restaurant?
 - What is the **Price**?
 - Is it **Raining**?
 - Do I have a **Reservation**?
 - What **Type** of food is it?
 - What's the **Estimated** wait time?
 - I will wait, **True** or **False**

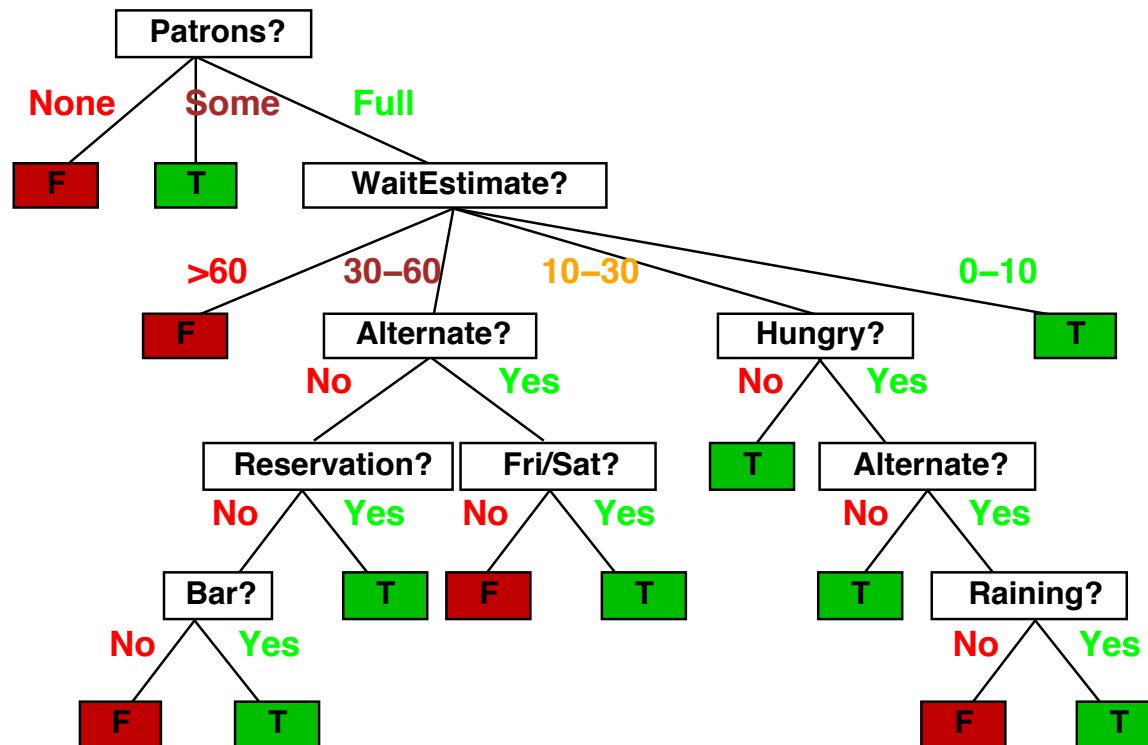
Attribute-Based Representations

- Items (or examples) described by attribute values (Boolean, discrete, continuous, etc.)
- Example: Will I wait for a table at a restaurant?

	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est	WillWait
X1	T	F	F	T	Some	\$\$\$	F	T	French	0–10	T
X2	T	F	F	T	Full	\$	F	F	Thai	30–60	F
X3	F	T	F	F	Some	\$	F	F	Burger	0–10	T
X4	T	F	T	T	Full	\$	F	F	Thai	10–30	T
X5	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
X6	F	T	F	T	Some	\$\$	T	T	Italian	0–10	T
X7	F	T	F	F	None	\$	T	F	Burger	0–10	F
X8	F	F	F	T	Some	\$\$	T	T	Thai	0–10	T
X9	F	T	T	F	Full	\$	T	F	Burger	>60	F
X10	T	T	T	T	Full	\$\$\$	F	T	Italian	10–30	F
X11	F	F	F	F	None	\$	F	F	Thai	0–10	F
X12	T	T	T	T	Full	\$	F	F	Burger	30–60	T


Decision Trees

- A Decision Tree asks questions about the attributes of a given example to provide an answer
- Here is the “true” tree for whether or not to wait:



Decision Trees

- Decision trees are expressive: they can express any function of the input attributes
- For example, for Boolean functions, we can convert *truth table row* \rightarrow *path to leaf*



- Trivially, there is a consistent decision tree for any training set w/ one path to leaf for each example (unless f nondeterministic in x)
- But it probably won't generalize to new examples!

Decision Trees

- How many distinct decision trees are there with n Boolean attributes?
 - = number of Boolean functions
 - = number of distinct truth tables with 2^n rows = 2^{2^n}
 - E.g., with 6 Boolean attributes, there are 18,446,744,073,709,551,616 trees
- How many purely conjunctive hypotheses are there? (e.g., Hungry $\wedge \neg$ Rain)
 - Each attribute can be positive, negative, or omitted
→ 3^n distinct conjunctive hypotheses
- A more expressive hypothesis space...
 - increases chance that target function can be expressed
 - increases number of hypotheses consistent w/ training set
 - But may get worse predictions!

Decision-Tree Learning

- **Goal:** find a small tree consistent with examples
- **Idea:** (recursively) choose the “most significant” attribute as the root of a (sub)tree

```
function DTL(examples, attributes, default) returns a decision tree
  if examples is empty then return default
  else if all examples have the same classification then return the classification
  else if attributes is empty then return MODE(examples)
  else
    best ← CHOOSE-ATTRIBUTE(attributes, examples)
    tree ← a new decision tree with root test best
    for each value  $v_i$  of best do
      examplesi ← {elements of examples with best =  $v_i$ }
      subtree ← DTL(examplesi, attributes − best, MODE(examples))
      add a branch to tree with label  $v_i$  and subtree subtree
  return tree
```


ID3

- Developed by Quinlan (1986)
- ID3 = “Iterative Dichotomiser 3”
 - Create a decision tree that iteratively evaluates an item based on its attributes
 - Key idea: use *entropy* (or *information gain*) to make these decisions

ID3

- Consider how different attributes might split a set of training examples
 - Based on **Patrons**...

	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est	WillWait
X1	T	F	F	T	Some	\$\$\$	F	T	French	0–10	T
X2	T	F	F	T	Full	\$	F	F	Thai	30–60	F
X3	F	T	F	F	Some	\$	F	F	Burger	0–10	T
X4	T	F	T	T	Full	\$	F	F	Thai	10–30	T
X5	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
X6	F	T	F	T	Some	\$\$	T	T	Italian	0–10	T
X7	F	T	F	F	None	\$	T	F	Burger	0–10	F
X8	F	F	F	T	Some	\$\$	T	T	Thai	0–10	T
X9	F	T	T	F	Full	\$	T	F	Burger	>60	F
X10	T	T	T	T	Full	\$\$\$	F	T	Italian	10–30	F
X11	F	F	F	F	None	\$	F	F	Thai	0–10	F
X12	T	T	T	T	Full	\$	F	F	Burger	30–60	T



ID3

- Consider how different attributes might split a set of training examples
 - Based on **Type**...

	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est	WillWait
X1	T	F	F	T	Some	\$\$\$	F	T	French	0–10	T
X2	T	F	F	T	Full	\$	F	F	Thai	30–60	F
X3	F	T	F	F	Some	\$	F	F	Burger	0–10	T
X4	T	F	T	T	Full	\$	F	F	Thai	10–30	T
X5	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
X6	F	T	F	T	Some	\$\$	T	T	Italian	0–10	T
X7	F	T	F	F	None	\$	T	F	Burger	0–10	F
X8	F	F	F	T	Some	\$\$	T	T	Thai	0–10	T
X9	F	T	T	F	Full	\$	T	F	Burger	>60	F
X10	T	T	T	T	Full	\$\$\$	F	T	Italian	10–30	F
X11	F	F	F	F	None	\$	F	F	Thai	0–10	F
X12	T	T	T	T	Full	\$	F	F	Burger	30–60	T

ID3



- Consider how different attributes might split a set of training examples



- Which of these attributes should we prefer?

ID3

- Consider how different attributes might split a set of training examples



- Which of these attributes should we prefer?
 - Better to have subsets that are all/mostly the same class
 - "Patrons"** is a better choice here

Entropy

- “Information” is what answers questions in the context of decision trees
 - In essence, the more clueless I am about the answer initially, the more information is contained in the answer
- Scale
 - 1 bit = answer to Boolean question with prior $\langle 0.5, 0.5 \rangle$
 - Information in an answer when prior is $\langle P_1, \dots, P_n \rangle$ is

$$H(\langle P_1, \dots, P_n \rangle) = \sum_{i=1}^n -P_i \log_2 P_i$$

(also called the **entropy** of the prior)

Entropy


- Given a set of symbols $S...$
- Drawn from an alphabet $B...$
- **Entropy** = expected number of bits needed to encode the next symbol (amount of information that knowing one more symbol provides us)
 - If S is drawn at random from B , entropy is maximal
 - If S is always the same symbol from B , entropy is minimal
 - we know which symbol will come next, so there's no new information)

Entropy

- Entropy can be defined as:

$$H(\langle P_1, \dots, P_n \rangle) = \sum_{i=1}^n -P_i \log_2 P_i$$

- Example for a binary variable:



ID3


- Suppose we have p positive and n negative examples at the root
 - $H(\langle \frac{p}{p+n}, \frac{n}{p+n} \rangle)$ bits needed to classify a new example
 - E.g., for the restaurant examples, $p = n = 6$, so 1 bit
- An attribute splits the examples E into subsets E_i , each of which (we hope) needs less information to complete the classification

ID3

- Let E_i have p_i positive and n_i negative examples
 - $H(\langle \frac{p_i}{p_i+n_i}, \frac{n_i}{p_i+n_i} \rangle)$ bits needed to classify a new example
 - the **expected** number of bits per example over all branches is

$$\sum_i \frac{p_i + n_i}{p + n} H(\langle \frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i} \rangle)$$

ID3




■ For **Patrons**:

- Subset E_1 has $p_1 = 0$ positive, $n_1 = 2$ negative
- Subset E_2 has $p_2 = 4$ positive, $n_2 = 0$ negative
- Subset E_3 has $p_3 = 2$ positive, $n_3 = 4$ negative

$$\sum_i \frac{p_i + n_i}{p + n} H \left(\left\langle \frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i} \right\rangle \right) = 0.459 \text{ bits}$$

ID3



- For **Type**:

- All subsets have $p_i = n_i$

$$\sum_i \frac{p_i + n_i}{p + n} H\left(\left\langle \frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i} \right\rangle\right) = 1 \text{ bit}$$

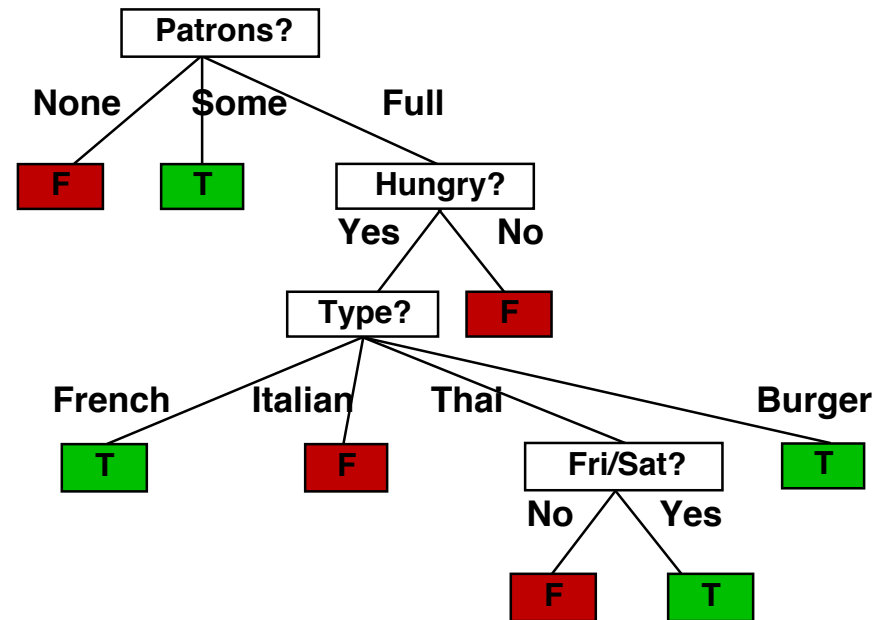
- So we choose the attribute that minimizes the remaining information needed, i.e., **Patrons**

ID3

- Alternatives to Entropy...
 - E.g., Information Gain is a common alternative
 - And there are others (Gain Ratio, GINI Index, RLDM, etc.)
- Alternative search strategies
 - E.g., Decision Forests (instead of 1 deep tree, use multiple shorter trees and aggregate their decisions)

ID3

- The ID3 tree after learning:



- A lot simpler than the “true” tree we saw earlier
- Why is simpler a good thing here?

Overfitting

- What if almost all examples fit a clear pattern, but there's one that doesn't fit?

	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est	WillWait
X1	T	F	F	T	Some	\$\$\$	F	T	French	0–10	F
X2	T	F	F	T	Full	\$	F	F	Thai	30–60	F
X3	F	T	F	F	Some	\$	F	F	Burger	0–10	F
X4	T	F	T	T	Full	\$	F	F	Thai	10–30	F
X5	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
X6	F	T	F	T	Some	\$\$	T	T	Italian	0–10	T
X7	F	T	F	F	None	\$	T	F	Burger	0–10	T
X8	F	F	F	T	Some	\$\$	T	T	Thai	0–10	T
X9	F	T	T	F	Full	\$	T	F	Burger	>60	T
X10	T	T	T	T	Full	\$\$\$	F	T	Italian	10–30	F
X11	F	F	F	F	None	\$	F	F	Thai	0–10	F
X12	T	T	T	T	Full	\$	F	F	Burger	30–60	F

Overfitting

- What if almost all examples fit a clear pattern, but there's one that doesn't fit?

	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est	WillWait
X1	T	F	F	T	Some	\$\$\$	F	T	French	0–10	F
X2	T	F	F	T	Full	\$	F	F	Thai	30–60	F
X3	F	T	F	F	Some	\$	F	F	Burger	0–10	T
X4	T	F	T	T	Full	\$	F	F	Thai	10–30	F
X5	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
X6	F	T	F	T	Some	\$\$	T	T	Italian	0–10	T
X7	F	T	F	F	None	\$	T	F	Burger	0–10	T
X8	F	F	F	T	Some	\$\$	T	T	Thai	0–10	T
X9	F	T	T	F	Full	\$	T	F	Burger	>60	T
X10	T	T	T	T	Full	\$\$\$	F	T	Italian	10–30	F
X11	F	F	F	F	None	\$	F	F	Thai	0–10	F
X12	T	T	T	T	Full	\$	F	F	Burger	30–60	F

Overfitting

- What happens if we have an example in the training set that is noise? (e.g. that is mislabeled)
 - ID3 will try to force it into the decision tree!
- This is an example of **Overfitting** — when what you've learned does very well on training examples, but much less well on other examples
 - In other words, it doesn't generalize
 - One of the most critical concepts in machine learning!
- For ID3, some strategies exist to avoid this
 - e.g., prevent leaves that have very few examples
 - e.g., stop building the tree when there would be little information gain in the next level

Overfitting

- More generally, what can you do?
 - Use a **training set** and a **test set**
 - Say, reserve 90% for training and 10% for testing
 - Use cross-validation
 - Repeatedly split your set into different training vs. test sets
 - E.g., Leave-one-out cross-validation
 - By not training on the data you're testing, you're trying to ensure generality

Next time...

- We will look at one particular type of Machine Learning called Reinforcement Learning...